

Diffusion Models

Diffusion Models



A hedgehog using a calculator.




A corgi wearing a red bowtie and a purple



A transparent sculpture-
a duck made out of glass.

High-level overview

- Diffusion models are probabilistic models used for image generation
- They involve reversing the process of gradually degrading the data
- Consist of two processes:
 - **The forward process:** data is progressively destroyed by adding noise across multiple time steps
 - **The reverse process:** using a neural network, noise is sequentially removed to obtain the original data



A high-level conceptual overview of the entire image space.

Three Categories

- Denoising Diffusion Probabilistic Models (DDPM)
- Noise Conditioned Score Networks (NCSN)
- Stochastic Differential Equations (SDE)

Notations

$p(x_0)$ - data distribution

$\mathcal{N}(x; \mu, \sigma \cdot I)$ - Gaussian distribution

Random Variable (image)

Mean Vector

Covariance matrix. I is the identity matrix

Forward Process (Iterative)


$$\beta_t \ll 1, t = \overline{1, T}$$

$$x_t \sim p(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t} \cdot x_{t-1}, \beta_t I)$$




x_0

...



x_{t-1}



x_t

...



x_T

Forward Process (One Shot)

$$x_t \sim p(x_t|x_0) = \mathcal{N}(x_t; \sqrt{\hat{\beta}_t} \cdot x_0, (1 - \hat{\beta}_t)I)$$

$$\hat{\beta}_t = \prod_{i=1}^t \alpha_i$$
$$\alpha_t = 1 - \beta_t$$



x_0

...



x_{t-1}



x_t

...



x_T




DDPMs. Properties of β_t

✓ 1. $\beta_t \ll 1, t = \overline{1, T}$

$$x_t \sim p(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t} \cdot x_{t-1}, \beta_t I)$$

x_t is created with a small step modeled by β_t




x_{t-1} comes from region close to x_t ,
therefore we can model with Gaussian

DDPMs. Properties of β_t

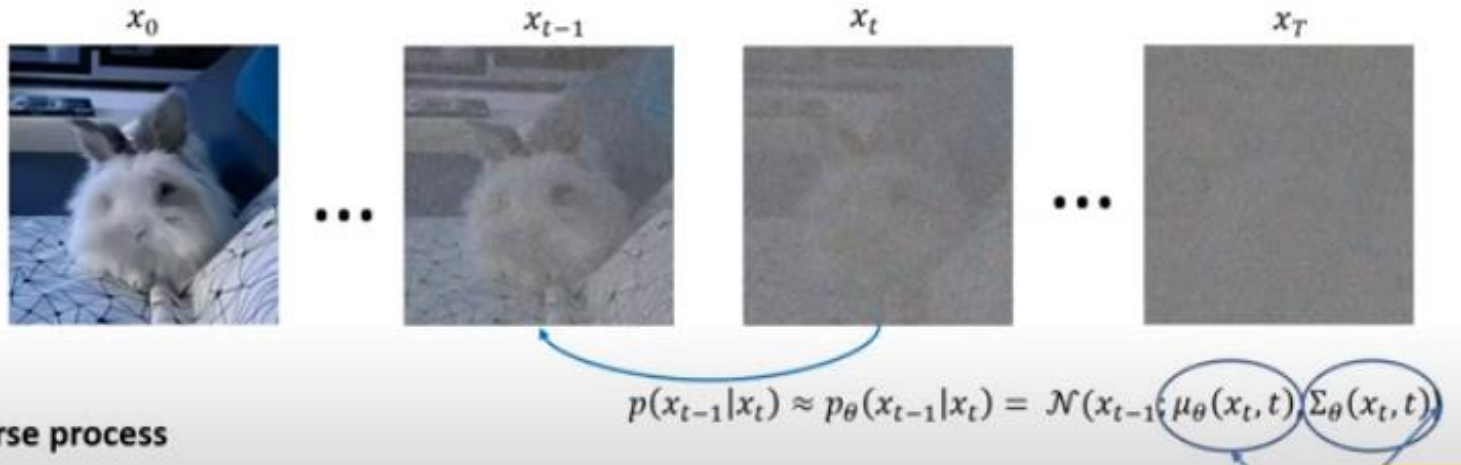
✗ 1. $\beta_t \ll 1, t = \overline{1, T}$

$$x_t \sim p(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t} \cdot x_{t-1}, \beta_t I)$$



DDPMs. Training objective

Remember that:



Neural
Network

Cross Entropy and KL (Kullback-Leibler) divergence

- Entropy: $E(P) = - \sum_i P(i) \log P(i)$
- Cross Entropy: $C(P) = - \sum_i P(i) \log Q(i)$
- KL divergence: $D_{KL}(P \parallel Q) = \sum_i P(i) \log [P(i)/Q(i)] = \sum_i P(i) [\log P(i) - \log Q(i)]$

DDPMs. Training Objective

$$\min_{\theta} \mathbb{E}_{x_0 \sim p(x_0)} \left[-\log p_{\theta}(x_0|x_1) + KL(p(x_T|x_0)||p_{\theta}(x_T)) + \sum_{t=2}^T KL(p(x_{t-1}|x_t, x_0)||p_{\theta}(x_{t-1}|x_t)) \right]$$

$$\min_{\theta} \mathbb{E}_{x_0 \sim p(x_0)} \left[-\log p_{\theta}(x_0|x_1) + \sum_{t=2}^T KL(p(x_{t-1}|x_t, x_0) || p_{\theta}(x_{t-1}|x_t)) \right]$$

Notations:

$$p(x_{t-1}|x_t, x_0) = \mathcal{N}(x_{t-1}; \tilde{\mu}(x_t, x_0), \tilde{\beta}_t I)$$

$$\tilde{\mu}(x_t, x_0) = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \hat{\beta}_t}} z_t \right), z_t \sim \mathcal{N}(0, I)$$

$$\tilde{\beta}_t = \frac{1 - \hat{\beta}_{t-1}}{1 - \hat{\beta}_t} \cdot \beta_t$$

$$\hat{\beta}_t = \prod_{i=1}^t \alpha_i$$

$$\alpha_t = 1 - \beta_t$$

$$\min_{\theta} \frac{1}{T} \sum_{t=1}^T \mathbb{E}_{x_0 \sim p(x_0), z_t \sim \mathcal{N}(0,1)} \|z_t - z_{\theta}(x_t, t)\|_2^2$$

DDPMs. Training Algorithm

$$\min_{\theta} \underbrace{\frac{1}{T} \sum_{t=1}^T \mathbb{E}_{x_0 \sim p(x_0), z_t \sim \mathcal{N}(0, I)} \|z_t - z_{\theta}(x_t, t)\|_2^2}_{\mathcal{L}_{simple}}$$

Training algorithm:

Repeat
 $x_0 \sim p(x_0)$
 $t \sim \mathcal{U}(\{1, \dots, T\})$
 $z_t \sim \mathcal{N}(0, I)$
 $x_t = \sqrt{\hat{\beta}_t} \cdot x_0 + \sqrt{(1 - \hat{\beta}_t)} z_t$
 $\theta = \theta - lr \cdot \nabla_{\theta} \mathcal{L}_{simple}$
Until convergence

$$\hat{\beta}_t = \prod_{i=1}^t \alpha_i$$

Algorithm 1 Training

1: **repeat**
2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
3: $t \sim \text{Uniform}(\{1, \dots, T\})$
4: $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
5: Take gradient descent step on
 $\nabla_{\theta} \left\| \boldsymbol{\epsilon} - \boldsymbol{\epsilon}_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}, t) \right\|^2$
6: **until** converged

Algorithm 2 Sampling

1: $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
2: **for** $t = T, \dots, 1$ **do**
3: $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ if $t > 1$, else $\mathbf{z} = \mathbf{0}$
4: $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon}_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$
5: **end for**
6: **return** \mathbf{x}_0

Thank You